

Distributed Systems, 5DV020, HT09  
Ludvig Widman(dit06lwn), Emil Eriksson(c07een)

Tutors: Lars Larsson, Daniel Henriksson  
Date: October 27, 2009  
Bonus level: 2

Running the program:

1. `$ cd ~dit06lwn/edu/DS/git/DistSys-Project/`
2. Start registry and sequencer: `$ ant seq`
3. Start a client: `$ ant gui`
4. Goto 3

## Assignment part 2

# GCom

# Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
<b>2</b>	<b>Tools and language for implementation</b>	<b>1</b>
<b>3</b>	<b>Usage guide: the demo application</b>	<b>1</b>
3.1	Starting the application . . . . .	1
3.2	Using the application . . . . .	1
3.2.1	Creating a group . . . . .	1
3.2.2	Joining a group . . . . .	2
3.2.3	Chatting in a group . . . . .	2
<b>4</b>	<b>System Overview</b>	<b>5</b>
4.1	The gcom package . . . . .	5
4.1.1	GCom . . . . .	5
4.1.2	BasicCommunicationModule . . . . .	6
4.1.3	ReliableCommunicationModule . . . . .	6
4.1.4	GroupManagementModule . . . . .	6
4.1.5	RMIModule . . . . .	6
4.1.6	momNonOrdered . . . . .	6
4.1.7	momFIFO . . . . .	6
4.1.8	momCausal . . . . .	6
4.1.9	momTotal . . . . .	7
4.1.10	ReferenceKeeper . . . . .	7
4.1.11	HashVectorClock . . . . .	7
4.1.12	GroupDefinition . . . . .	7
4.1.13	RemoteObject . . . . .	7
4.1.14	Group . . . . .	7
4.1.15	Member . . . . .	7
4.1.16	Message . . . . .	7
4.1.17	Debug . . . . .	7
4.2	The rmi package . . . . .	7
4.2.1	RMIServer . . . . .	8
4.2.2	BackdoorOpener . . . . .	8
4.2.3	Sequencer . . . . .	8
4.2.4	SequencerCommunicationModule . . . . .	8
4.2.5	GroupSequencer . . . . .	8
4.3	The gui package . . . . .	8
4.3.1	GUIViewOther . . . . .	8
4.3.2	GroupPanel . . . . .	8
<b>5</b>	<b>Message ordering and message ordering modules</b>	<b>9</b>
5.1	Non ordered / momNonOrdered . . . . .	9
5.2	FIFO ordering / momFIFO . . . . .	9
5.3	Causal ordering / momCausal . . . . .	10
5.4	Total ordering / momTotal . . . . .	11
5.5	Causal total ordering . . . . .	11
<b>6</b>	<b>Design decisions</b>	<b>13</b>
6.1	Selection of bonus level . . . . .	13
6.2	Our implementation of static groups . . . . .	13
6.3	Why an chat Application? . . . . .	13
6.4	Why using a sequencer for total ordering? . . . . .	13
6.5	Why the sequencer lives in the rmi-server . . . . .	14

6.6	Why use a backdoor? . . . . .	14
6.7	Is a rmi registry neccessary? . . . . .	14
6.8	All clients are assumed to be good . . . . .	14
6.9	Reliable multicast sends some extra messages to itself . . . . .	14
<b>7</b>	<b>Fault tolerance</b>	<b>14</b>
7.1	Detecting crashed members . . . . .	14
7.2	If the registry restarts or is unresponsive . . . . .	15
7.3	If the group leader crashes . . . . .	15
7.3.1	Election of a new leader . . . . .	15
7.4	If the last member of a group crashes . . . . .	15
7.5	Netsplit . . . . .	16
<b>8</b>	<b>Realization of the project plan</b>	<b>16</b>
<b>9</b>	<b>Implementation of requirement specification</b>	<b>16</b>
<b>10</b>	<b>Limitations</b>	<b>16</b>
<b>A</b>	<b>Project plan</b>	<b>17</b>
<b>B</b>	<b>Requirement specification from project plan</b>	<b>17</b>
B.1	Group management . . . . .	17
B.2	Communication . . . . .	18
B.3	Message Ordering . . . . .	18

## 1 Introduction

In this project we have created a middleware for group communication. A demo application to demonstrate that the middleware works has also been created.

The middleware handles communication between members of groups. Groups can have different message orderings, they can be static or dynamic and they can use either reliable or basic multicast for sending messages. Every member of the group uses the same of these properties.

Each group has a leader which has a reference in a registry. Clients can find groups and join them by getting a group list from the registry and contacting the leader. The middleware can be used independently of the demo application but this will not be described in this report.

Communication in groups is done by passing messages. A message is usually sent to the entire group, only when joining a group messages is sent to only one member in the group. Messages have a type so that application messages can be separated from system messages.

The middleware can be used independently from the demo application but this will not be described further in this report.

## 2 Tools and language for implementation

Since the clear recommendation was to use Java and Java RMI we have decided to follow that recommendation. While there are languages that would be more interesting and easier to design GUIs in (cocoa/objective c), Java is more easily ported and the available help from lecturers also helped tip the balance in Java's favor.

JUnit has be used for unit testing and logging has be done through the log4j-package.

## 3 Usage guide: the demo application

Before you can use the application you must obtain the source code. The code is available at CS servers at the following location: `~dit06lwn/edu/DS/git/DistSys-Project/`

### 3.1 Starting the application

Before starting the client, start a rmi regisry and a sequencer on a known host with:

```
$ ant seq
```

The client can be started from the command line with:

```
$ ant gui
```

### 3.2 Using the application

When the client is started you must begin with connecting to a registry. This can be done from the File menu and then Connect to registry. A dialog like the one in figure 1 on the following page is displayed. Enter the hostname and port of the registry and click connect.

If the client is connected to a registry you can create or join groups.

#### 3.2.1 Creating a group

To create a group, select Create group from the Group menu. You are now presented with the dialog in figure 3 on page 3. Enter the name of the group and your




Figure 1: The dialog for connecting to a rmi registry.

Ordering	Explanation
NONORDERED	Messages are not ordered at all.
FIFO	Messages are ordered per sender in the order they where sent.
CAUSAL	Messages are ordered in a “happen before” relation.
TOTAL	Messages are ordered in the exact same order by all clients in the group
CAUSALTOTAL	Messages are ordered in a causal order and in the exact same order by all clients

Figure 2: Breif explanations of the different message orderings.

nickname.

Select group type from the dropdown box, in dynamic groups users can join and leave at anytime, static groups has a fixed set of members.

Choose which type of multicast you would like to use in the group. Basic multicast just sends the messages, reliable multicast guarantees that an attempt to send the message to all users has been made before the message is delivered to the sender.

Finally, select type of ordering and click Create. Please refer to figure 2 for descriptions of the available message ordering types.

### 3.2.2 Joining a group

To join a group, select Join group from the Group menu. A dialog like the one in figure 4 on the next page is displayed. Enter a nickname and select a group from the Groups list. Do not attempt to join backdoor or sequencer. Click Join.

### 3.2.3 Chatting in a group

When you have joined a group you will see the group in the tab bar in the main window. Click the groups name to select it (see figure 5 on page 4). When a group is selected messages from the group is displayd to the left. To the right a list of members in the group is shown. At the bottom there is a textfield for entering messages to the group. You send a message by pressing Enter in the textfield.

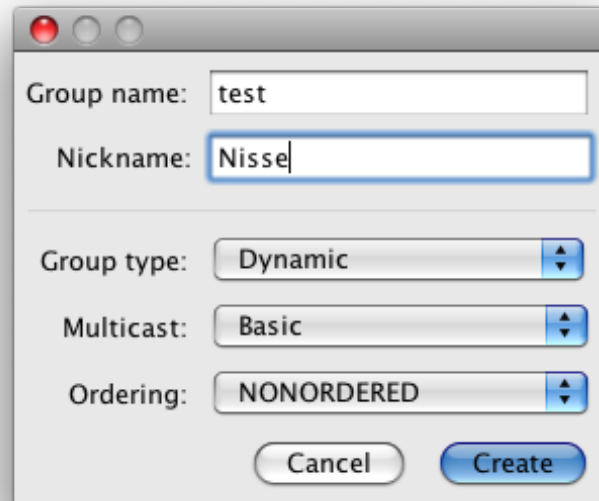


Figure 3: The dialog for creating a new group.




Figure 4: The dialog for joining a group.




Figure 5: A chat room the user has joined that contains no other users.

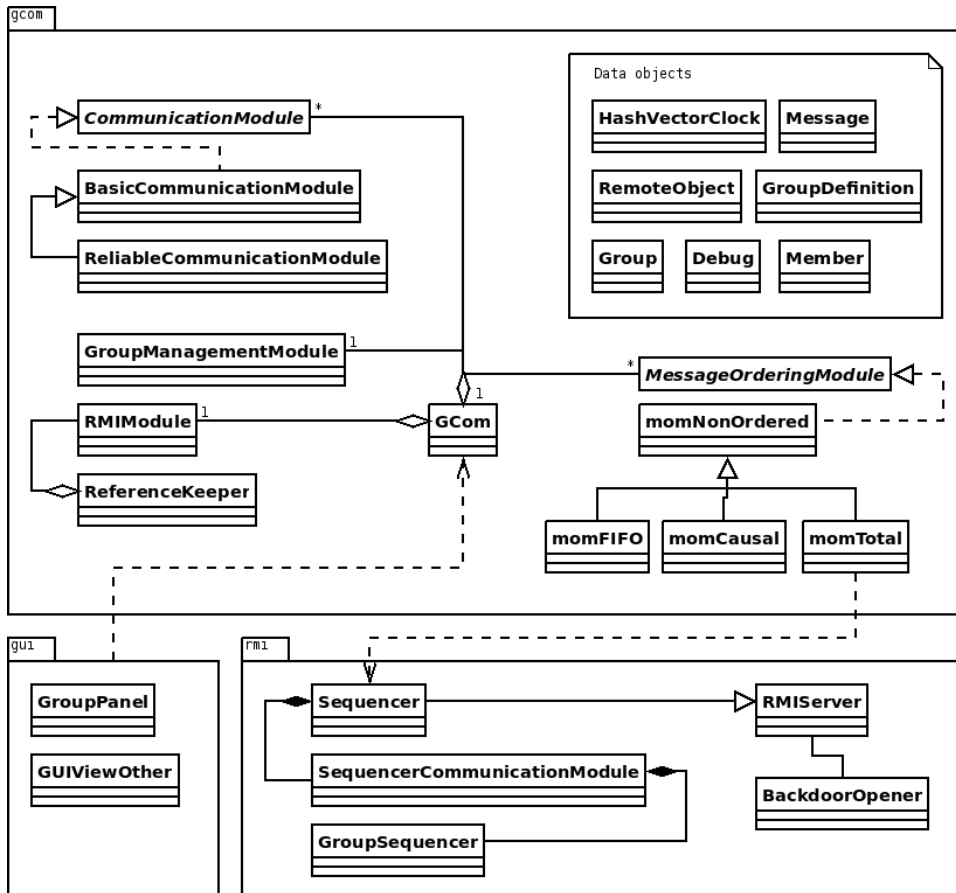


Figure 6: All the classes in the program and an overview of their relations.

## 4 System Overview

This project consists of a middleware for group communication and a demo application to show that it works. The middleware is located in the gcom and rmi packages, the demo application is in the gui package. An overview over all the packages and their classes can be seen in figure 6.

### 4.1 The gcom package

The gcom package is the main part of this project. This is where communication, message ordering and group management is handled.

#### 4.1.1 GCom

GCom is the main component in this package and is where the demo application is connected. The application gives commands to GCom, for example “create a group” or “send a message”, these commands are translated to messages that GCom sends through the appropriate communication module and to commands to the GroupManagementModule so that the current group state can be maintained.

Information also flows from the communication modules via the message ordering modules back to GCom when messages from other clients is received. These messages either effects the group state or is application data that GCom forwards to the application.



GCom has one GroupManagementModule, one RMIModule and the same amount of communication and message ordering modules as the number of groups it handles.

#### 4.1.2 BasicCommunicationModule

The BasicCommunicationModule (BCM) is the simplest implementation of the interface CommunicationModule. It implements basic multicast and has the ability to send and receive messages. Each communication module corresponds to a specific group.

When a message is sent the BCM sends it to every member of the group, including the sender. It also detects if the connection is refused to any of the senders and removes these members from the group via a message to GCom.

When the BasicCommunicationModule receives a message the message is sent directly to the message ordering module for the group.

#### 4.1.3 ReliableCommunicationModule

The ReliableCommunicationModule (RCM) extends the BCM and implements reliable multicast.

When the RCM receives a message it checks if it has received it before. If not the message is added to a cache of received messages. The message is then sent to all group members via BCM and delivered to a message ordering module.

Sending messages in the RCM also stores the message in the cache. The message is then sent to all members of the group and the message is delivered to the sender.

When a message is added to the cache, the size of the cache is also trimmed by removing the oldest message if the cache is larger than a specified max size.

#### 4.1.4 GroupManagementModule

The GroupManagementModule (GMM) stores information about groups and can notify listeners when something changes.

The main responsibilities for the GMM includes storing group state, resolving group names to members and updating group state when changes are received.

#### 4.1.5 RMIModule

The RMIModule connects to a rmi registry at a given host and port. It has an interface for binding, unbinding, listing and rebinding objects.

#### 4.1.6 momNonOrdered

The momNonOrdered is the simplest implementation of the MessageOrderingModule (MOM) interface. Every group has a corresponding message ordering module. The type depends on what the creator of the group chose.

Message ordering is described in detail in section 5 on page 9.

#### 4.1.7 momFIFO

The momFIFO orders messages in FIFO ordering, in which messages is delivered in the same order the sender sent them. See section 5.2 on page 9 for details.

#### 4.1.8 momCausal

Causal ordering orders messages according to a “happened before” relationship. See section 5.3 on page 10 for details.

#### **4.1.9 momTotal**

Total ordering means that all clients deliver the messages in the exact same order. See section 5.4 on page 11 for details.

#### **4.1.10 ReferenceKeeper**

The ReferenceKeeper is used only by GCom for the group leader. It checks that the rmi registry has the reference to the leaders RemoteObject and tries to rebind it if not. The ReferenceKeeper is run in it's own thread and rechecks the reference in the registry every minute.

#### **4.1.11 HashVectorClock**

The HashVectorClock is a simple vector clock. It stores a clock state with integer values for different process IDs. The clock can be compared to another clock, merged with other clocks and ticked, which increases the clock value for the owner of the clock by one.

#### **4.1.12 GroupDefinition**

A GroupDefinition stores information about a group like message ordering, communication type and if the group is static or dynamic.

#### **4.1.13 RemoteObject**

The RemoteObject is used for communication between processes via java RMI.

The object has a queue in which all incoming messages are queued. A thread is responsible for transferring messages from the queue to the corresponding CommunicationModule.

#### **4.1.14 Group**

Group stores members and GroupDefinition for a specific group. This class is only used by the GroupManagementModule for storing state.

#### **4.1.15 Member**

Member stores id, name and RemoteObject for a member in a group.

#### **4.1.16 Message**

Message is a serializable class used for sending information between processes. A message contains a vector clock, group name, source, message type and message data. Messages used for requesting serial numbers from the sequencer also has a return address.

#### **4.1.17 Debug**

Debug holds a log4j logger and is used for printing debug and error messages.

### **4.2 The rmi package**

The rmi package contains classes related to the rmi registry.

#### 4.2.1 RMIServer

The RMI server creates a registry and opens a backdoor for binding objects remotely.

#### 4.2.2 BackdoorOpener

BackdoorOpener is used by RMIServer to create a backdoor. The backdoor is needed since the used implementation of the registry does not allow remote processes to bind objects.

#### 4.2.3 Sequencer

The Sequencer is used for getting a serial number on messages that all clients can agree to. There is only one sequencer per group, and in our implementation only one in total.

Sequencer extends the RMI-server and therefore first starts a registry. It then binds in a SequencerCommunicationModule SCM as “sequencer” in the registry which the clients that uses any of the total orderings will connect to.

Clients can request serial numbers by sending a message to the sequencer and supplying a return address in the message. The Sequencer will order the messages it receives in either a momNonOrdered or a momCausal depending on which of total or causal total ordering is requested. Messages that has already been sequenced are cached and the same serial number will be returned if a new request for the same message arrives.

#### 4.2.4 SequencerCommunicationModule

The SequencerCommunicationModule (SCM) implements a CommunicationModule since this makes communication with it convenient, but the module itself is only used by the Sequencer.

It receives messages and returns them to the return address with a serial number. The serial numbers is supplied by a GroupSequencer. Each group has it's own GroupSequencer.

#### 4.2.5 GroupSequencer

GroupSequencer (GS) is used by SCM for setting serial numbers on messages. The GS uses a MOM to order the messages before a serial number is given to them. Messages are cached with their serial numbers so requests for the same message gets the same number.

### 4.3 The gui package

The gui package implements our demo application: a chat application with tabs for groups.

#### 4.3.1 GUIViewOther

A simple chat client implemented with GCom.

#### 4.3.2 GroupPanel

A class implementing a panel containing a textarea, nicklist and textfield. One is created for each group.

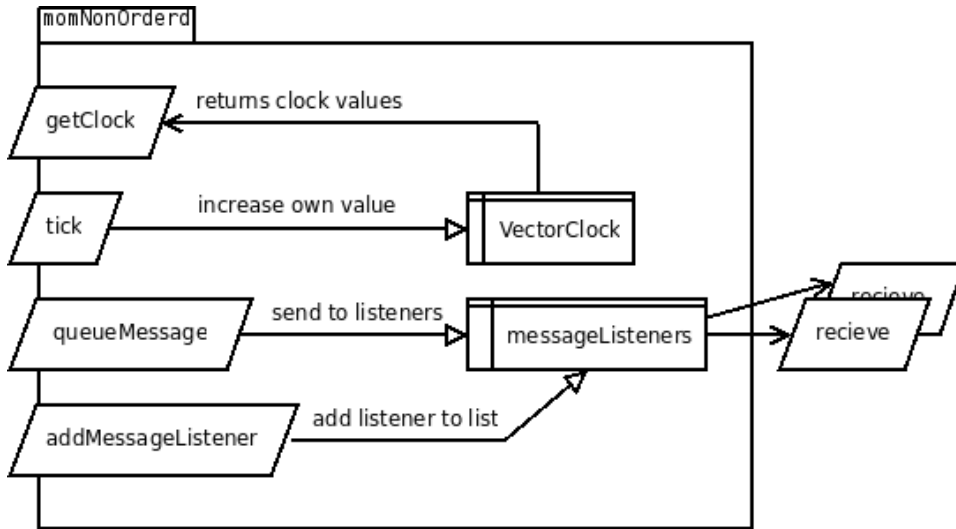


Figure 7: The non ordered message ordering module just forwards messages to it's listeners. The state of a vector clock is also kept in the module.

## 5 Message ordering and message ordering modules

GCom features five kinds of message orderings: NonOrdered, FIFO, Causal, Total and CausalTotal. The first four of these has their own message ordering module. Each type of ordering and it's corresponding class is described in detail below.

The message ordering modules have each been thoroughly tested with unit tests to make sure that the algorithms works like intended.

Additional information about algorithms and definitions of the ordering types above can also be found in [1].

### 5.1 Non ordered / momNonOrdered

The non ordered ordering module (MOM) does not order messages at all. This is the base class for the ordering modules. Every message ordering module has a queuing method, a vector clock and message listeners. All the NonOrdered does is to receive messages in the queue method and send them through to it's messages listeners directly. This is also shown in figure 7.

This module also has two functions related to the vector clock. It increases the vector clock when told so and returns the clock values if asked. These functions is used in several of the child modules that inherits from momNonOrdered.

### 5.2 FIFO ordering / momFIFO

The FIFO ordered MOM delivers messages in the order they where sent. This works by numbering every sent message with a vector clock and storing received messages in a hold back queue until they can be delivered in order. Messages in the queue is compared to the previous last delivered message from the same sender. Messages from new senders is delivered directly as no clock value is known. This is also described in figure 8 on the next page.

In detail the sorting algorithm works as follows:

1. Check if no previous messages has been received from this sender, if so:
  - (a) Send message to message listeners.

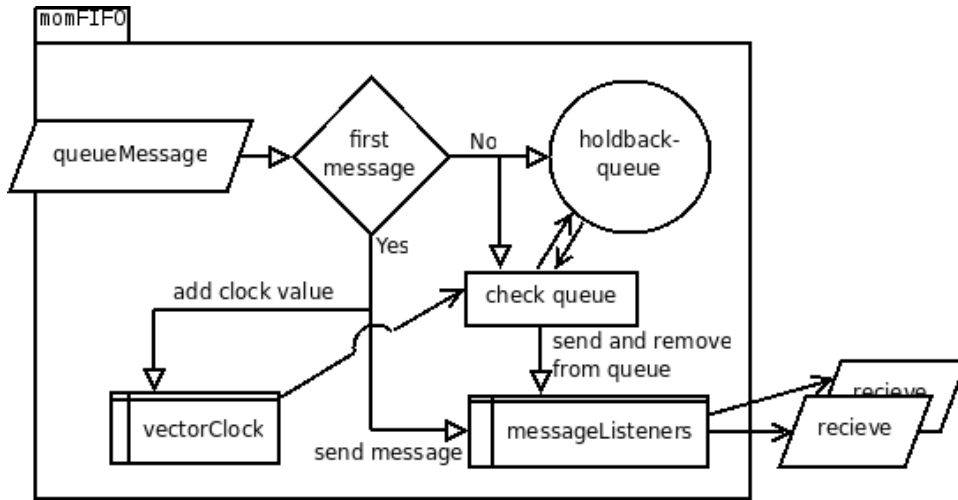


Figure 8: The FIFO ordered MOM places messages in a hold back queue until they can be delivered in the order they where sent. Messages from new senders are delivered directly as no previous clock value is known.

- (b) Increase local clock value for sender by one.
2. If not, add message to the hold back queue.
3. Loop through the hold back queue and with any message that has a clock value that equals the local clock value plus one:
  - (a) Send message to message listeners.
  - (b) Increase local clock value for sender by one.
  - (c) Remove message from hold back queue.

### 5.3 Causal ordering / momCausal

Causal ordering is an ordering where a message can not be delivered before a message that happened before it. Just like in the FIFO ordering the message is first placed in a hold back queue. If a new sender is detected their clock value is added to the local clock.

Messages in the queue is checked when a message is received. If all previous messages from the same sender has been delivered, and all messages delivered by the sender when the message was timestamped has been delivered, then the message can be delivered. See figure 9 on the following page.

In detail the sorting algorithm works as follows:

1. Check if no previous messages has been received from this sender, if so:  
Add senders clock value minus one, from message to local clock
2. If not, add message to the hold back queue.
3. Loop through the hold back queue and with any message that:
  - (a) Has a clock value for the sender that equals the local clock value for the sender plus one
  - (b) Has clock values before or equal to corresponding local clock values, excluding the value for the sender

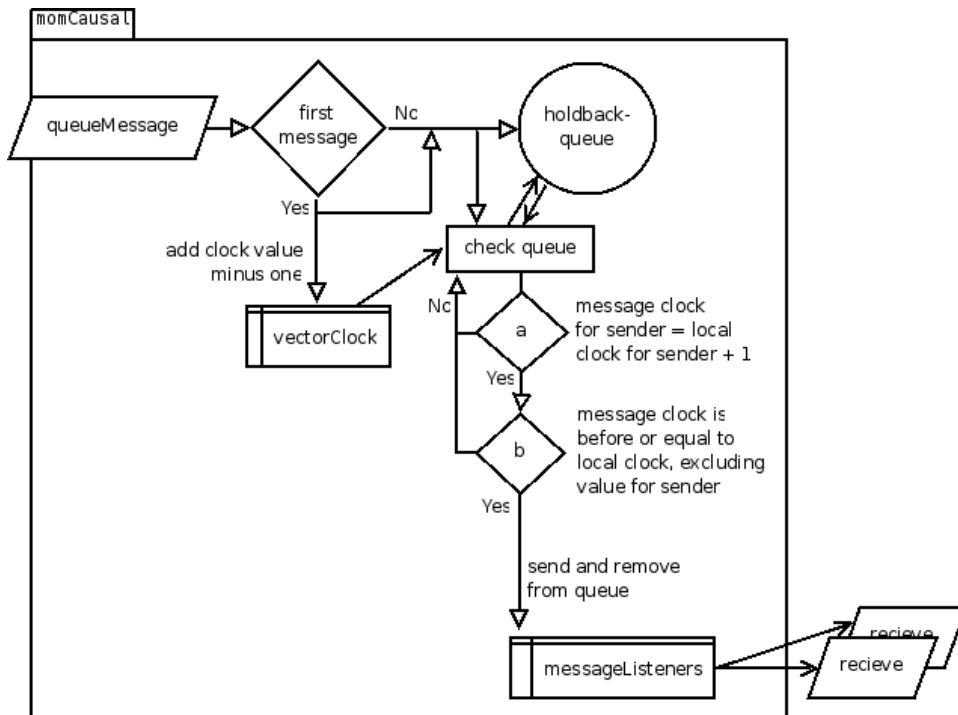


Figure 9: The Causal ordered MOM places a message in a hold back queue until all messages that happened before it has been delivered.

Do the following:

- (a) Send message to message listeners.
- (b) Increase local clock value for sender by one.
- (c) Remove message from hold back queue.

In step 1 the clock value minus one is added to the local clock, this is since the message can not be delivered directly like in FIFO ordering (the message might not conform to criteria b).

#### 5.4 Total ordering / momTotal

Total ordering is an ordering where every client delivers the messages in the same order. In our implementation this is solved by the use of a sequencer, which is an application that is global for the group.

When a message is received it is sent to the sequencer. The sequencer gives the message a serial number and sends it back. When the MOM gets a message with a serial number it puts it in hold back queue until it can be delivered in sequence (until all messages with a lower number is delivered). This is also described in figure 10 on the next page.

The first message with a serial number the MOM receives is delivered directly and any message with a lower number than that is discarded.

#### 5.5 Causal total ordering

Causal total ordering uses the same MOM as total ordering. The only difference is in the sequencer.

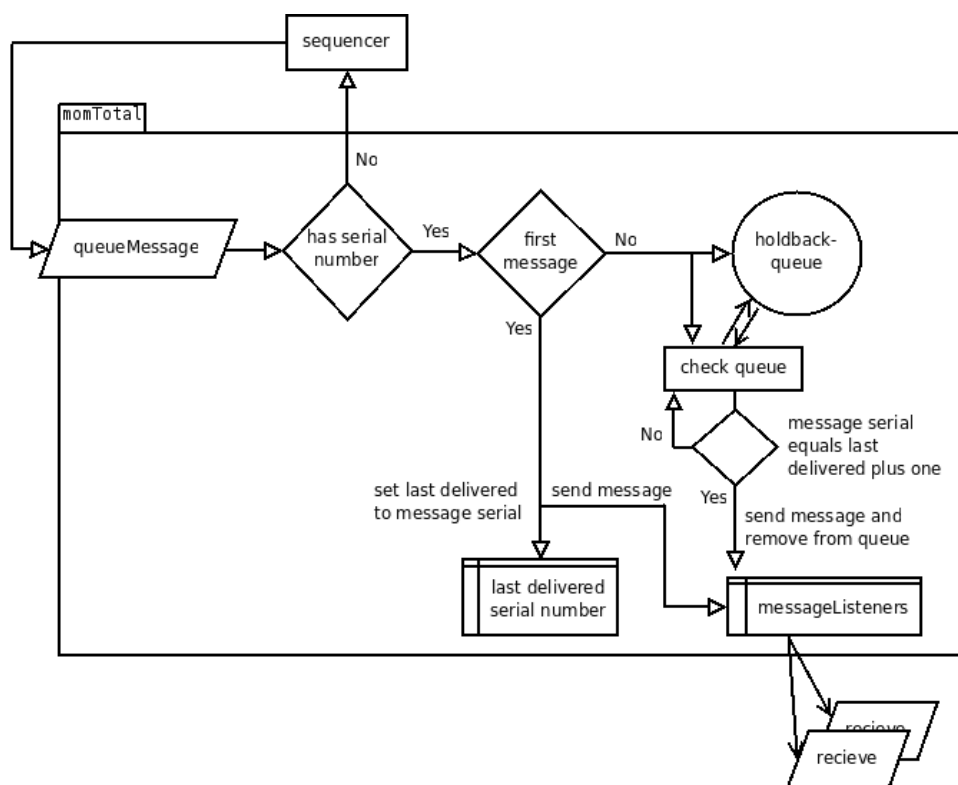


Figure 10: Total ordering uses a sequencer to deliver messages in the same order in every client.




Figure 11: Causal total ordering also uses a sequencer, but here the sequencer has a MOM of it's own to give the messages serial numbers in causal order.

With causal total ordering the sequencer uses a momCausal to order the messages in causal order and gives them serial numbers in the order they are delivered. This interaction is shown in figure 11.

## 6 Design decisions

### 6.1 Selection of bonus level

We chose to solve bonus level two with dynamic groups since it seemed like a reasonable trade off between extra credits and extra workload.

### 6.2 Our implementation of static groups

We have implemented static groups as dynamic groups with some extra limitations. The list of members in a static group is created dynamically (by letting users join freely), the group is then frozen and the group becomes static. This was the easiest way to implement static groups when we already had the dynamic ones. If only static groups would have been implemented they probably could have been created in a simpler way.

### 6.3 Why a chat Application?

As our demo application on top of GCom we built a chat application with tabs for different groups. This type of application maps well to the underlying structure of members that belongs to groups. It also gave an easy way to see that messages arrive unharmed and in the correct amounts and order.

### 6.4 Why using a sequencer for total ordering?

When implementing total ordering one can either use a sequencer or a distributed algorithm. The sequencer method has the downside of the sequencer being the single point of failure. However we still choose that method since it seemed more straightforward to implement and we already had another single point of failure.



## 6.5 Why the sequencer lives in the rmi-server

The RMI-server has to be started before the sequencer and we only wanted a single instance of both so it made sense to put them in the same module (or rather, make the sequencer extend the RMI-server). The RMI-server is also run on a hopefully quite stable host which makes it a good location for the sequencer.

It would have scaled to more users to put the sequencer in the group leader, but this would make the sequencer less stable (since the group leader might choose to leave the group).

An obvious drawback of having a single sequencer for all groups is that it becomes a bottleneck. This could be solved by moving the sequencer to several dedicated hosts and assigning groups to one of many sequencers.

## 6.6 Why use a backdoor?

We could either have implemented our own registry, have one registry per host, or create a backdoor for binding objects from remote hosts. The latter seemed like a easier solution.

## 6.7 Is a rmi registry necessary?

During the analysis phase we speculated in if the registry really would be necessary. We actually sketched a more distributed design where a client only needed to find any node in the network to get a list of groups. This would be solved by letting nodes ask other nodes for known groups and leaders. This solution was quite early deemed unpractical and much too complicated.

## 6.8 All clients are assumed to be good

We are assuming that all clients are good and that no one is lying. This simplifies things significantly and it seemed a bit out of scope to ensure that everything works correctly against all possible kinds of attacks. The security in the system is generally very low. If arbitrary messages can be sent (for example by writing a new application ontop of gcom) a user can do almost anything.

## 6.9 Reliable multicast sends some extra messages to itself

Due to laziness and because it does not really affect anything reliable multicast sends messages to itself twice and discards one of the messages. It also resends unseen received messages to itself and discards the resended message. This made the implementation simpler and it hardly takes any extra time.

# 7 Fault tolerance

## 7.1 Detecting crashed members

The communication module detects if a client refuses connections or if connections timeout. If this happens the member is removed from the group management module and a “lost member” message is sent to the group.

If the crashed client is the group leader a new leader is elected.




Figure 12: Two members engaged in election.

## 7.2 If the registry restarts or is unresponsive

The leader of each group occasionally checks that the group is registered in the RMI-registry. If the registry has restarted for some reason the leader will re-register the group.

## 7.3 If the group leader crashes

If the group leader crashes the group will detect that no messages arrive to the leader and then elect a new leader.

### 7.3.1 Election of a new leader

When a member (A in figure 12) detects that the leader is missing, an election of a new leader will take place. The member sets itself as leader of the group and sends an election message containing the members unique id to all known members of the group.

Member (B) receives an election message. If the member hasn't already begun an election it sends its id to all members of the group and sets either itself or the sender as leader depending on which has the higher id. If the member has already begun the election process the received id is compared to the saved largest id. The larger of the two is stored and the sending member is set as leader if the sent id is higher.

When a member has received election messages from all other members, the election is done and if a member has set itself as leader of the group it creates a ReferenceKeeper and binds a reference in the registry.

This algorithm isn't foolproof since crashing members can cause other members to wait for responses indefinitely.

## 7.4 If the last member of a group crashes

If the last member of a group crashes no one will unregister the group. This will potentially lead to a lot of "dead" references in the registry if there is a bug causing nodes to crash when they are the last to leave a group. In order to counter this extra testing has to be done to make sure that nodes behave when leaving and unregistering with the registry.

There is currently no way to clear out old references in the registry (other than restarting it periodically).

## 7.5 Netsplit

If the group is split in two partitions, due to a cable failure or similar, the partition with the leader will continue to work as before. The other partition will elect a new leader when they discover that they no longer receive messages from the leader.

The new leader will try to reregister the group in the registry. The leader for the other part will do the same and therefore compete over the reference in the registry. Because of this newcomers will join the group having the reference in the registry at the time the node connects.

## 8 Realization of the project plan

The project plan from deliverable 1 is included in appendix A on the next page. We have been slightly behind the plan during most of the project part due to the new influenza that struck Emil during the first weeks.

We also moved forward the design of the GUI to milestone 4 since it was easier to do while implementing the actual application. We also moved total and causal total ordering forward one week since the other message orderings took most of the assigned week.

Generally things took slightly longer then anticipated, part due to tricky bugs and part due to tricky algorithms.

## 9 Implementation of requirement specification

The requirement specification from the project plan is included in appendix B on the following page. All the requirements of the specification are implemented.

## 10 Limitations

- The RMI-registry is a single point of failure. If the registry crashes no processes can find or join groups until the registry has been restarted.
- All processes are assumed to be good. If a process starts to lie or send malicious messages bad things will happen.
- If a client crashes during election undefined things happen.
- If the sequencer crashes all groups using the sequencer must be restarted since the message cache is lost and messages will get new serial numbers starting at 1. The sequencer is placed at the same host as the registry which is assumed to be stable enough for this to not be a big problem.
- If a group contains too many users the java VM will run out of memory and the client will crash. If reliable multicast is used the network will most likely also be very congested.
- The demo application does not use the remove group feature of GCom since it is more user friendly to let the group live as long as there is users connected to it and remove it when the last user quits. However, the feature is implemented and can be used by some other application that uses the middleware.

- The program does not build under Windows. Usage of windows is discourage and behaviour is undefined.
- In the demo application backdoor and sequencer is visible in the groups list.
- If too many totally ordered messages is sent (probably a bit over 500000 messages with the default VM memory size) the sequencer will crash since it does not empty it's cache. This is easy to fix, but we realised this a bit too late to have the time to implement the fix.

## References

- [1] Distributed systems : concepts and design / GeorgeCoulouris, Jean Dollimore, Tim Kondberg.-4th ed. p. 490-498

## A Project plan

Date	Milestone	Content
11 Sept	Part 1 due	Written project plan
18 Sept	Milestone 1	Create all interfaces Build dataobjects Design GUI mockups for test app
25 Sept	Milestone 2	Implement basic methods like non-ordered ordering and non-reliable multicast.
2 Oct	Milestone 3	Implement advanced methods like total ordering and reliable multicast.
9 Oct	Milestone 4	Implement GUI-application and debug mode.
16 Oct	Milestone 5	Resolve distributed problems. Write on the report. Make test protocol.
23 Oct	Finalization	Everything should be finished. Minor adjustments.
27 Oct	Hand in report	
28-30 Oct	Demonstration	

## B Requirement specification from project plan

### B.1 Group management

**REQUIRED** External processes must be able to find and connect to group leaders via a RMI registry.

**REQUIRED** The group leader must regularly make sure that the group and leader is registered in the RMI registry.

**REQUIRED** Any process must be able to create a new group and register itself as leader in the registry.

**REQUIRED** Any process must be able to join a group by connecting to the group leader.

**REQUIRED** Every process in a group must notify the other processes of any known changes to the group composition.

**REQUIRED** Every process must keep track of the members of all the groups it's a member of.

**REQUIRED** Every process must monitor if another process in the group stops responding and report a corresponding change in group composition.

**REQUIRED** If the leader of the group leaves or crashes a new leader must be elected in the group.

## B.2 Communication

**REQUIRED** It must be possible to send messages to the members of joined group.

**REQUIRED** Both non-reliable and reliable multicast must be supported for sending messages.

**REQUIRED** The communication method chosen by the group creator (reliable or non-reliable multicast) must be used when communicating in a group.

**REQUIRED** Every message must contain a updated vector clock.

## B.3 Message Ordering

**REQUIRED** Any message must arrive in proper order (according to the selected sorting requirements).

**REQUIRED** A message must not be delivered until all previous messages are delivered if the message ordering demands so.